


Picking ONE compute-efficient HP-search recipe for training student models

(LegNet students distilled from a frozen AlphaGenome oracle; goal = strong, diverse model ensemble at every dataset size D)



Why one deep run per cell? A strategy's value IS the shape of its accuracy-vs-compute curve, so 'optimal compute' (the knee) and 'best strategy' (the height) are read from the SAME curves.